

Step 1: Verify the API

Step 2: Login to cloudshell

Step 3: Create a subnet in default VPC

```
gcloud compute networks subnets create gke-deep-dive-subnet --network=default --range=10.10.0.0/24
```

Step 4: Verify the created network and it's subnet

**No Secondary IP ranges**

Step 5: Create VPC native cluster

```
gcloud container clusters create gke-deep-dive --num-nodes=1 --disk-type=pd-standard  
--disk-size=10 --enable-ip-alias --subnetwork=gke-deep-dive-subnet --addons=HttpLoadBalancing
```

Step 6: Verify the secondary IP ranges created

Step 7: Verify the cluster is deployed okay

Step 8: In the console, check the VPC-native traffic routing under networking section of the GKE cluster - it should be enabled

**The new cluster should have the [HttpLoadBalancing](#) add-on enabled, Check**

**If not, enable it by running the command below:**

```
gcloud container clusters update gke-deep-dive --update-addons=HttpLoadBalancing=ENABLED
```

## Step 9: Create necessary files:

### 1. gke-deep-dive-app.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: gke-deep-dive-app
spec:
  replicas: 2
  selector:
    matchLabels:
      app: online
  template:
    metadata:
      labels:
        app: online
    spec:
      containers:
        - image: gcr.io/google-containers/echoserver:1.10
          name: gke-deep-dive-app
          ports:
            - name: http
              containerPort: 8080
          readinessProbe:
            httpGet:
              path: /healthz
              port: 8080
              scheme: HTTP
```

### 2. gke-deep-dive-svc.yaml:

```
apiVersion: v1
kind: Service
metadata:
  name: gke-deep-dive-svc
  annotations:
    cloud.google.com/l4-rbs: "enabled"
spec:
  type: LoadBalancer
  externalTrafficPolicy: Cluster
  selector:
    app: online
  ports:
    - name: tcp-port
      protocol: TCP
      port: 8080
      targetPort: 8080
```

Please notice the option: cloud.google.com/l4-rbs: "enabled"

It instructs GKE to create a backend service-based external passthrough Network Load Balancer **so that clients outside the cluster can send packets to the Service's Pods.**

Step 10: Create the deployment:

```
kubectl apply -f gke-deep-dive-app.yaml
```

Step 11: Verify that there are two serving Pods for the Deployment:

```
kubectl get pods
```

Step 12: Create the service:

```
kubectl apply -f gke-deep-dive-svc.yaml
```

Step 13: Verify that your Service is running:

```
kubectl get svc gke-deep-dive-svc
```

**GKE assigned an *EXTERNAL\_IP* for the external passthrough Network Load Balancer.**

**Note:** It might take a few minutes for GKE to allocate an external IP address before the load balancer is ready to serve your application.

**Test connecting to the load balancer:**

```
curl EXTERNAL_IP:PORT
```

Step 14: Verify the external LoadBalancer Service and its components

```
kubectl describe svc gke-deep-dive-svc
```

**Step 15: Delete the *gke-demo-svc* external LoadBalancer Service.**

```
kubectl delete service store-v1-lb-svc
```

Step 16: Delete the cluster

```
Gcloud container clusters delete gke-deep-dive
```