

We'll create a private cluster named `gke-deep-dive` that has private nodes, and has no client access to the public endpoint. As part of the same command, we'll also create

- A subnet named `gke-deep-dive-subnet`

Step 1: `gcloud config set compute/zone us-west1-a`

Step 2: `gcloud container clusters create gke-deep-dive --num-nodes=1`

`--disk-type=pd-standard --disk-size=10 --create-subnetwork name=gke-deep-dive-subnet  
--enable-ip-alias --enable-private-nodes --enable-private-endpoint --master-ipv4-cidr  
172.16.0.32/28`

where:

- `--create-subnetwork name=gke-deep-dive-subnet` causes GKE to automatically create a subnet named `gke-deep-dive-subnet`.
- `--enable-ip-alias` makes the cluster VPC-native
- `--enable-private-nodes` indicates that the cluster's nodes do not have external IP addresses.
- `--enable-private-endpoint` indicates that the cluster is managed using the private IP address of the control plane API endpoint.
- `--master-ipv4-cidr 172.16.0.32/28` specifies an internal IP address range for the control plane (optional for Autopilot).

Step 4: Verify that the cluster's nodes do not have external IP addresses.

`gcloud container clusters describe gke-deep-dive`

Step 5: Check the new subnet that's created on console

At this point, these are the only IP address ranges that have access to the control plane:

- The primary range of `gke-deep-dive-subnet`.
- The secondary range used for Pods.

Let's try to access the control plane from outside `gke-deep-dive-subnet`, using

**cloud shell. We must authorize the public IP address of our cloud shell instance to have access to the cluster endpoint.**

Step 6: `dig +short myip.opendns.com @resolver1.opendns.com`

Step 7: `gcloud container clusters describe gke-deep-dive --zone us-west1-a --format "flattened(masterAuthorizedNetworksConfig.cidrBlocks[])"`

Step 8: `gcloud container clusters update gke-deep-dive --zone us-west1-a --enable-master-authorized-networks --master-authorized-networks <IP address from step 6>`

Step 9: Error because this cluster doesn't have public endpoint

Create another cluster without private endpoint, limited access to public endpoint

Step 10: `gcloud container clusters create gke-deep-dive-public --num-nodes=1 --disk-type=pd-standard --disk-size=10 --enable-master-authorized-networks --subnetwork gke-deep-dive-subnet --enable-private-nodes --enable-ip-alias --master-ipv4-cidr 172.16.0.16/28`

**At this point, these are the only IP addresses that have access to the cluster control plane:**

- The primary range of `gke-deep-dive-subnet`.
- The secondary range used for Pods.

**Let's try to access the control plane from outside `gke-deep-dive-subnet`, using cloud shell. We must authorize the public IP address of our cloud shell instance to have access to the cluster endpoint.**

Step 11: gcloud container clusters describe gke-deep-dive-public --zone us-west1-a --format

```
"flattened(masterAuthorizedNetworksConfig.cidrBlocks[])"
```

Step 12: gcloud container clusters update gke-deep-dive-public --zone us-west1-a --enable-master-authorized-networks

--master-authorized-networks <IP address from step 6>

Step 13: gcloud container clusters describe gke-deep-dive-public --zone us-west1-a --format

```
"flattened(masterAuthorizedNetworksConfig.cidrBlocks[])"
```

**Now these are the only IP addresses that have access to the control plane:**

- The primary range of **gke-deep-dive-subnet**.
- The secondary range used for Pods.
- Address ranges that we have authorized for the cloud shell public IP

*Step 14: gcloud container clusters get-credentials gke-deep-dive-public --internal-ip*

Step 15: Kubectl get nodes

***It might take sometime***

*Step 16: If connection error, try: gcloud container clusters update gke-deep-dive-public --enable-master-global-access*

*Step 17: Delete cluster: gcloud container clusters delete -q gke-deep-dive-public gke-deep-dive --zone us-west1-a*